

# Green Model Advisor - Technology Guide

---

*Understanding the Technologies Behind Smart AI Model Selection*

---

## Table of Contents

1. [What is Green Model Advisor?](#)
  2. [Backend Technologies](#)
  3. [Frontend Technologies](#)
  4. [AI & LLM Technologies](#)
  5. [Carbon Tracking Technology](#)
  6. [How Everything Works Together](#)
  7. [The Technology Flow](#)
- 

## What is Green Model Advisor?

**In Simple Terms:** A smart system that helps you choose the best AI model for your task while tracking how much electricity and carbon it uses. Think of it as a "GPS for AI" - it finds the most efficient route (model) for your needs.

### The Problem It Solves:

- Multiple AI models exist (Google's, Mistral's, NVIDIA's)
- Each has different quality and power consumption
- Hard to compare manually
- No tracking of environmental impact

### The Solution:

- Automatically tests all models
  - Compares quality (accuracy) vs environmental cost (carbon)
  - Picks the best one for your task
  - Shows you the environmental impact
- 

## Backend Technologies

### 1. FastAPI - The Server Brain

#### What is it?

- A modern Python web framework that handles requests
- Like a restaurant's kitchen that receives orders and prepares responses

#### What does it do in our project?

- Receives your prompts from the website

- Coordinates which AI model to use
- Tracks carbon emissions
- Saves results to database
- Sends results back to your browser

## Why FastAPI?

- Super fast (modern Python)
- Auto-generates documentation
- Easy to understand code
- Handles many requests simultaneously

**Real-world analogy:** FastAPI = Restaurant manager who takes your order, tells the kitchen what to cook, and brings your plate back

---

## 2. SQLAlchemy + SQLite - The Data Vault

### What are they?

Technology	Role	Analogy
<b>SQLAlchemy</b>	Talks to database	Librarian
<b>SQLite</b>	Stores data	Library

### What they do:

- Store every request, response, and carbon measurement
- Keep history of what models were used
- Remember previous conversations

### Where is data stored?

```
green-model-advisor/
└── carbon_estimates.db ← All data here (SQLite file)
```

### What data is saved?

- Your prompt (input text)
- Which model was used
- Response generated
- Carbon emissions (kg CO<sub>2</sub>)
- Energy consumed (kWh)
- Accuracy score (0-100%)
- Timestamp

### Why SQLite?

- Lightweight (no server needed)

- Fast for single computer
  - Perfect for learning/demo projects
  - Easy to backup (just copy the file)
- 

### 3. Pydantic - The Data Validator

#### What is it?

- A Python library that checks data is correct before processing

#### How it works:

```
User Input → Pydantic checks → Is it valid?  
↓  
If YES: Process it  
If NO: Send error back
```

**Real-world analogy:** Pydantic = Airport security checking your ticket is real before letting you board

#### In our project:

- Validates your prompt isn't empty
  - Checks model name is valid
  - Ensures parameters make sense
- 

## ⌚ Frontend Technologies

### 1. React - The User Interface

#### What is it?

- A JavaScript library for building interactive websites
- Used by Facebook, Netflix, Instagram

#### What it creates:

- The chatbox where you type
- The buttons to select models
- The dashboard showing statistics
- The comparison results

#### How it works (simplified):

1. User types in textbox
2. React detects the change
3. Updates what you see on screen instantly
4. When you click "Submit", it sends to backend

5. Backend responds with results
6. React updates screen with results

## Why React?

- Updates screen without reloading page
  - Reusable components (buttons, cards, etc.)
  - Very fast and responsive
  - Huge community support
- 

## 2. Vite - The Build Tool

### What is it?

- A tool that packages your React code for the browser
- Like a compiler for web applications

### What it does:

- Combines all React files into one fast file
- Starts a development server (your local testing)
- Makes code run super fast

### Development server:

```
Terminal: npm run dev
          ↓
http://localhost:5173 ← Your app runs here
```

## Why Vite?

- Lightning fast (10x faster than older tools)
  - Modern and popular
  - Great for development
- 

## 3. Vanilla CSS - The Styling

### What is it?

- Plain CSS (no libraries) for colors, fonts, layout

### What it makes:

- Green and blue color scheme
- Nice buttons and cards
- Responsive design (works on phone/tablet/computer)
- Dark mode support

**File location:**

```
client/src/ChatCarbon.css ← All styling here
```

---

## 🤖 AI & LLM Technologies

### What is an LLM?

#### LLM = Large Language Model

Think of it as a very smart autocomplete:

- You type a prompt
  - It predicts what you want
  - Generates a complete response
  - Like typing on your phone, but way smarter
- 

## 1. Google Gemini 2.0 Flash

### What is it?

- Google's latest AI model
- Specialist in reasoning and understanding

### How to use it:

- Need API key from Google Cloud
- Sends your prompt to Google's server
- Gets response back in seconds

### Characteristics:

- Most accurate for code generation (77% accuracy)
- Balanced carbon emissions (~67 µg CO<sub>2</sub>)
- ⏳ Medium speed
- 💸 Uses Google's API

### Use case:


- Writing complex code
  - Explaining difficult concepts
  - General-purpose AI tasks
- 

## 2. Mistral 7B (via OpenRouter)

### What is it?

- Open-source AI model from Mistral AI
- Accessed through OpenRouter API

#### How it works:



#### Characteristics:

- Most carbon efficient (~60-65 µg CO<sub>2</sub>)
- Fast response
- Good accuracy (75-77%)
- Cheaper than Gemini
- Open-source (code available)

#### Use case:

- Budget-conscious tasks
- When you care more about speed
- Environmental consciousness


---

### 3. NVIDIA Qwen 2.5 Coder

#### What is it?

- Specialized model from NVIDIA
- Built specifically for CODE generation
- Uses NVIDIA's data center GPUs

#### How it works:



**Characteristics:**

- 🔧 Specialized for coding
- ✅ Good accuracy (70-73%)
- ⚡ Uses GPU acceleration (fast)
- 🏢 Enterprise-grade reliability
- 🔒 Secure API

**Use case:**

- Writing code (main purpose)
- Enterprise deployments
- Production systems

---

## How Model Selection Works

**The Algorithm (simplified):**

```
Step 1: Analyze Prompt
└── Is it a coding task?
└── How complex is it?
└── How long is it?
└── What domain (web/app/data)?  
  
Step 2: Test Each Model
└── Gemini Test
    ├── Carbon: 67 µg
    ├── Accuracy: 77%
    └── Score: 0.23
└── Mistral Test
    ├── Carbon: 62 µg
    ├── Accuracy: 76%
    └── Score: 0.24
└── NVIDIA Test
    ├── Carbon: 67 µg
    ├── Accuracy: 72%
    └── Score: 0.28  
  
Step 3: Choose Best
└── LOWEST SCORE WINS! (Mistral with 0.24)
```

**Scoring Formula:**

$$\text{Score} = \text{Carbon} + (1 - \text{Accuracy}/100)$$

Lower = Better

- Balances environmental impact
- Considers quality
- Fair comparison

## 🌐 Carbon Tracking Technology

### CodeCarbon - The Emission Tracker

#### What is it?

- Python library that measures real CPU/GPU power consumption
- Like a utility meter for your code

#### How it works:

1. Start measurement  
↓
2. Run AI model  
↓
3. Measure: CPU power, GPU power, Time duration  
↓
4. Calculate: Power × Time = Energy  
↓
5. Convert: Energy × Grid Carbon Intensity = CO<sub>2</sub>  
↓
6. Report: X grams of CO<sub>2</sub> emitted

#### Real numbers example:

```
Running Gemini for 2 seconds
└── CPU Power: 100 watts
└── GPU Power: 50 watts
└── Total: 150 watts × 2 seconds = 300 joules
└── Grid Intensity: 0.22 kg CO2 per kWh (US average)
└── Result: ~0.067 mg CO2 (67 micrograms)
```

#### Why this matters:

- AI models run on electricity
- Electricity comes from power plants
- Power plants emit CO<sub>2</sub> (mostly)
- More electricity = More CO<sub>2</sub> = More climate impact

#### What CodeCarbon measures:

- CPU usage
- GPU usage (if available)
- Memory usage
- Time duration

- Regional grid carbon intensity
- Total CO<sub>2</sub> emissions

### Accuracy:

- Real measurements (not estimates)
  - Varies by region (US vs Europe vs Asia have different grids)
  - Accounts for actual hardware used
- 

## NumPy - The Math Library

### What is it?

- Fast mathematical operations library

### In our project:

- Calculates average accuracy
  - Combines scores
  - Does statistical analysis
  - Works with carbon numbers
- 

## ⌚ How Everything Works Together

### Complete Request Flow

- ```
YOU (Browser)
  ↓
1. Type prompt: "Generate Python bubble sort code"
2. Click "Estimate" button
  ↓
REACT (Frontend)
  ↓
3. Collects your input
4. Shows "Processing..." message
5. Sends request to backend
  ↓
FASTAPI (Backend)
  ↓
6. Receives request
7. Validates with Pydantic
  ↓
MODEL SELECTION
  ↓
8. Tests all 3 models:
  a) Call Google Gemini API
  b) Call OpenRouter/Mistral API
  c) Call NVIDIA NIM/Qwen API
  ↓
CODECARBON (Each request)
```

- ↓
  9. Measures carbon emissions for each  
↓

CODE QUALITY EVALUATOR

↓

  - 10. Scores each response (0-100)  
↓

COMPARISON

↓

  - 11. Calculates scores for each model
  - 12. Picks best (lowest score)  
↓

SQLALCHEMY & SQLITE

↓

  - 13. Saves everything to database:
    - Prompt
    - All 3 responses
    - Carbon for each
    - Accuracy for each
    - Best model choice

↓

FASTAPI Response

↓

  - 14. Sends back best result  
↓

REACT (Frontend)

↓

  - 15. Receives data
  - 16. Updates screen with:
    - Carbon: 67 µg CO<sub>2</sub>
    - Energy: 0.22 J
    - Tokens: 245
    - Accuracy: 77%
    - Response: "def bubble\_sort..."

↓

YOU (Browser)

↓


  - 17. See results with nice formatting
  - 18. Can copy the code
  - 19. See which model was best

---

## #[ The Technology Stack Visualization

YOUR BROWSER

- REACT (Interactive UI)
- ChatCarbon.jsx (Main chat interface)
  - Dashboard.jsx (Statistics and metrics)
  - Vanilla CSS (Beautiful styling)



## 💡 The Technology Flow (Step-by-Step)

### Phase 1: Frontend Interaction

```
You (User)
↓
Type: "Write Python code for binary search"
↓
React detects input change
↓
Shows live character count (React state)
↓
You click "Estimate"
↓
React validates input (not empty?)
↓
Shows "Processing..." (React loading state)
```

### Phase 2: Backend Processing

```
FastAPI receives request
↓
Pydantic validates data
  |- Is prompt a string?
  |- Is model_name valid?
  |- Are parameters reasonable?
  ↓ (Yes to all)
Create unique request ID (UUID)
↓
Start timer (for tracking time)
```

### Phase 3: Multi-Model Testing

```
FOR EACH of 3 models:
  |- Model 1: Gemini
    |- Send request to Google API
    |- Receive response
    |- Start CodeCarbon tracker
    |- Process response (tokens, etc.)
    |- Stop CodeCarbon tracker
    |- Get emissions (67 µg CO2)
    |- Evaluate code quality → 77%
    |- Calculate score: 0.23
  |- Model 2: Mistral
    |- Send request to OpenRouter API
```

```

    └── Receive response
    └── Start CodeCarbon tracker
    └── Process response
    └── Stop CodeCarbon tracker
    └── Get emissions (62 µg CO2)
    └── Evaluate code quality → 76%
    └── Calculate score: 0.24 ← BEST!

└── Model 3: NVIDIA
    └── Send request to NVIDIA NIM
    └── Receive response
    └── Start CodeCarbon tracker
    └── Process response
    └── Stop CodeCarbon tracker
    └── Get emissions (67 µg CO2)
    └── Evaluate code quality → 72%
    └── Calculate score: 0.28

```

## Phase 4: Data Storage

```

SQLAlchemy creates database record:
    └── ID: 42
    └── prompt: "Write Python code for binary search"
    └── model_name: "mistral-7b-openrouter" ← Winner!
    └── carbon_emitted_kgco2: 0.000062 (62 µg)
    └── energy_consumed_kwh: 0.00019 (0.19 J)
    └── overall_accuracy: 76 (%)
    └── response_text: "def binary_search(arr, target)...)"
    └── total_tokens: 245
    └── created_at: 2025-12-01 14:32:15
    └── estimation_method: "codecarbon"
    ↓
SQLite saves to file: carbon_estimates.db

```

## Phase 5: Response to Frontend

```

FastAPI sends JSON response:
{
    "id": 42,
    "model_name": "mistral-7b-openrouter",
    "carbon_emitted_kgco2": 0.000062,
    "energy_consumed_kwh": 0.00019,
    "overall_accuracy": 76,
    "response_text": "def binary_search(arr, target)...)",
    "tokens_input": 8,
    "tokens_output": 237,
    "total_tokens": 245,
    "created_at": "2025-12-01T14:32:15Z"
}

```

## Phase 6: Frontend Display

```

React receives response
↓
Updates state with results
↓
Shows metrics in nice boxes:
  |- Carbon: 62 µg CO2
  |- Energy: 0.19 J
  |- Accuracy: 76%
  |- Tokens: 245
↓
Shows generated code in code block
↓
Adds copy button
↓
Hides "Processing..." message
↓
You see beautiful formatted results!

```

## Key Technologies Explained for Beginners

### Why These Technologies?

| Tech              | Why We Chose It    | Benefits                             |
|-------------------|--------------------|--------------------------------------|
| <b>React</b>      | Popular & powerful | Smooth, fast UI updates              |
| <b>FastAPI</b>    | Modern Python      | Quick development, auto-docs         |
| <b>SQLite</b>     | Simple & reliable  | No setup needed                      |
| <b>CodeCarbon</b> | Real measurements  | Actual CO <sub>2</sub> , not guesses |
| <b>Three LLMs</b> | Best diversity     | Accuracy + speed + price options     |

## What Students Learn From This Project

### 1. Full-Stack Development

- Frontend: React, CSS, JavaScript
- Backend: Python, FastAPI, APIs
- Database: SQL, ORM

### 2. API Integration

- REST APIs
- Multiple API providers

- Error handling
- Rate limiting

### 3. AI/ML Concepts

- LLM fundamentals
- Model comparison
- Scoring algorithms
- Accuracy evaluation

### 4. Environmental Tech

- Carbon tracking
- Energy measurement
- Sustainability
- Green computing

### 5. Software Architecture

- Separation of concerns
  - Client-server model
  - Data flow
  - System design
- 

## ✿ Real-World Applications

### Where This Technology is Used:

#### 1. Enterprise AI Platforms

- Companies choose between multiple AI providers
- Cost and performance optimization
- Carbon-aware infrastructure

#### 2. Research Institutions

- Studying AI model efficiency
- Carbon footprint of AI
- Environmental impact research

#### 3. Cloud Services

- AWS, Google Cloud, Azure
- Recommending best model/service
- Cost optimization

#### 4. Sustainability Companies

- Tracking carbon emissions
- Making green choices
- Corporate sustainability reports

## 5. Educational Platforms

- Teaching full-stack development
  - Understanding AI models
  - Learning system design
- 

## Performance Characteristics

### Response Times

```
API Request: ~50-500 ms (depending on model)
CodeCarbon Measurement: ~2-5 seconds per request
Database Save: ~50-100 ms
Frontend Update: <100 ms
Total: ~3-6 seconds per estimate
```

### Resource Usage

```
Memory:
└ Backend: ~200-300 MB
└ Database: Grows with each request
└ Frontend: ~100-150 MB

Disk:
└ Code: ~50 MB
└ Dependencies: ~500 MB
└ Database: Starts at <1 MB, grows with usage

Network:
└ Per request: ~10-50 KB upload
└ Per response: ~5-20 KB download
```

## Security Features

### Data Protection

- Input validation (Pydantic)
- CORS enabled (safe API access)
- No API keys in code (use .env)
- Database encryption ready

### Best Practices

- Environment variables for secrets
- HTTPS ready (production)

- Rate limiting ready
  - Error handling
- 

## Future Technology Enhancements

### 1. More LLM Models

- OpenAI GPT
- Claude
- Local LLMs (Llama)

### 2. Advanced Analytics

- Machine Learning for predictions
- Trend analysis
- Pattern recognition

### 3. Scalability

- PostgreSQL (bigger database)
- Redis (caching)
- Kubernetes (container orchestration)

### 4. Real-time Features

- WebSockets (live updates)
- Streaming responses
- Real-time dashboards

### 5. Mobile Support

- React Native app
  - Mobile-optimized interface
  - Progressive Web App (PWA)
- 

## Learning Resources

### Technologies Covered

- Python web development
- JavaScript React
- REST APIs
- Database design
- Cloud API integration
- Environmental monitoring

### Skills Gained

- Full-stack development
- System design

- API integration
  - Data analysis
  - Sustainability awareness
- 

## ⌚ Key Takeaways

### **Green Model Advisor combines:**

- Modern web technologies
- AI/ML services
- Real-time monitoring
- Environmental consciousness

### **It demonstrates:**

- Full-stack architecture
- Multi-service integration
- Practical AI application
- Sustainability in tech

### **Perfect for:**

- Learning full-stack development
  - Understanding AI models
  - Exploring green computing
  - Building production systems
- 

## 📞 Questions About Technologies?

### **Want to know more about a specific technology?**

- Check the main README.md for setup
  - Look at source code comments
  - Review the API documentation at [/docs](#)
  - Check technology official websites:
    - React: [react.dev](https://react.dev)
    - FastAPI: [fastapi.tiangolo.com](https://fastapi.tiangolo.com)
    - CodeCarbon: [codecarbon.io](https://codecarbon.io)
    - SQLAlchemy: [sqlalchemy.org](https://sqlalchemy.org)
- 

Made with  to help you understand modern AI technology